An interpretation of Bayesian global optimization as proposed by Snyman and Fatti [1]

Dirk Munro (Hamburg, Germany)

January 3, 2022

1 Confidence-based global optimization

Assume that we have an optimization problem and a search algorithm (e.g. gradient-based) to solve it. That is, for a given starting position $\mathbf{x}_0$ the algorithm will return a stationary point $\mathbf{x}_*$, a candidate optimum from all the potential optima—in general there are plenty of local optima, and one global optimum.

Assume that the optimization problem has a finite but unknown number of optima $\mathbf{x}_*^j$, with $j = 1, 2, \ldots, s$, and that from a particular starting point $\mathbf{x}_0$, the algorithm will converge to one of them. This procedure which transforms a given $\mathbf{x}_0$ to one of the the optima $\mathbf{x}_*^j$ is deterministic, yet unpredictable. In other words, given a new $\mathbf{x}_0$ (pulled 'from a hat', so to speak), we do not know *a priori* which $\mathbf{x}_*^j$ will be returned; however, given the same $\mathbf{x}_0$ again, the same $\mathbf{x}_*^j$ (from before) will be returned.

Lets say an an 'experiment' consists of sampling a random starting point $\mathbf{x}_0$ from the design domain, and running the optimization algorithm. There is hence a probability P_*^j that a particular $\mathbf{x}_*^j$ will be returned. Also, there is a probability P_* that the (global) optimum will be returned.

What is the probability P_* of finding the global optimum from a random starting point? If we had an estimate of the total number of optima s, then a reasonable estimate may be $1/s$...? Or rather, the flexible form of the Beta distribution may be assumed

$$\mathbf{p}[\mathrm{P}_*] = \frac{1}{\beta\,[a,b]} {\mathrm{P}_*}^{a-1}(1-\mathrm{P}_*)^{b-1}\,, \tag{1}$$

which reflects the *probability density function* $\mathbf{p}$ of the probability P_* of finding the global optimum from a random starting position—see for example Figure 1. An expected value of P_* given an a and b, and other statistical measures, is implied by (1).

Bayes' theorem permits us to update the prior probability density function $\mathbf{p}\,[\mathrm{P}_*]$, given the results from a number of experiments. Assume (for now) that we know the number of 'successes' r, after a total number of experiments n. That is, running the algorithm from n random starting positions, we assume that we can measure the number of times r the global optimum was found (r 'successes')[1]. Given n and r, the prior probability

[1]The severity/subtlety of this assumption—*i.e.*, knowing which is the global optimum, and hence being able to measure the amount of times it is found—is addressed later.

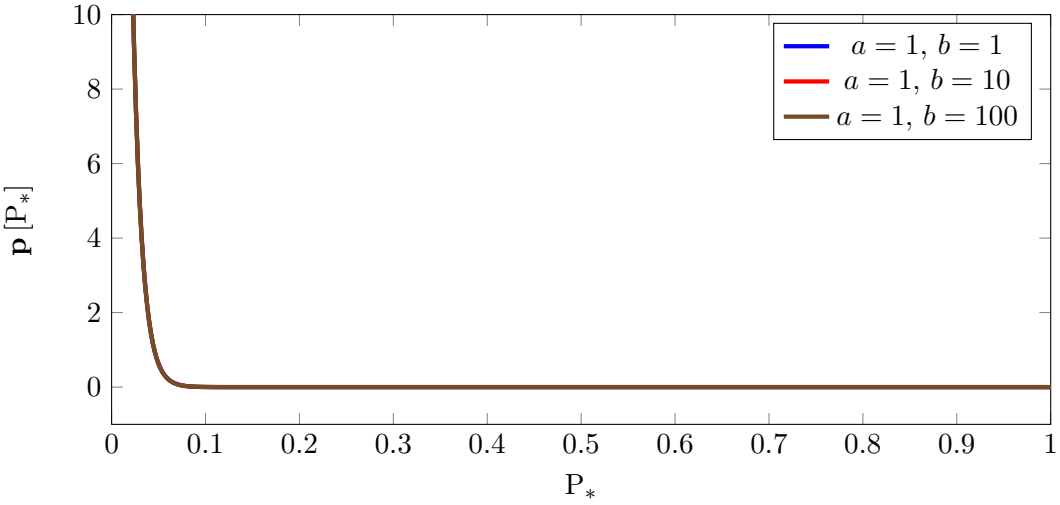

Figure 1: Example Beta distributions (probability density functions) of P_*, for representative values of a and b.

density function of P_* may be updated as per Bayes

$$\mathbf{p}[\mathrm{P}_* \,|\, [n, r]] = \mathbf{p}\,[r \,|\, [n, \mathrm{P}_*]]\,\mathbf{p}\,[\mathrm{P}_*] \Big/ \int_0^1 \mathbf{p}\,[r \,|\, [n, \mathrm{P}_*]]\,\mathbf{p}\,[\mathrm{P}_*]\,\mathrm{d}\mathrm{P}_* \tag{2}$$

The new term $\mathbf{p}\,[r \,|\, [n, \mathrm{P}_*]]$ is the probability density function of the number of 'successes' r, given n experiments and a prior probability P_*. This probability density function is thus of the binomial form

$$\mathbf{p}\,[r \,|\, [\mathrm{P}_*, n]] = \frac{n!}{r!(n-r)!} {\mathrm{P}_*}^r (1 - \mathrm{P}_*)^{n-r}\,. \tag{3}$$

In the end, we want to have a measure of confidence in whether we may have found the (global) optimum of the optimization problem, given n experiments, r 'successes', and the prior P_*. On the one hand, this can be expressed as the probability of having found the true optimum *at least once*, given n and P_*

$$\mathrm{P}\,[r_* \geq 1 |\, [n, \mathrm{P}_*]] = 1 - (1 - \mathrm{P}_*)^n\,, \tag{4}$$

which follows from the binomial distribution (3) of the complementary probability: 1 minus the probability of not finding the true optimum at all ($r = 0$, after n experiments). On the other hand, we have a probability density function $\mathbf{p}\,[\mathrm{P}_* |\, [n, r]]$, updated with the results r of n experiments, as per Bayes' theorem (2). Applying a change-of-variable [2] to $\mathbf{p}\,[P_* |\, \ldots]$ from Eq. (2), with Eq. (4) the transformation function, the probability of having found the optimum at least once, in terms of n and r, is

$$\mathrm{P}\,[r_* \geq 1 |\, [n, r]] = \int_0^1 \mathrm{P}\,[r^* \geq 1 |\, [n, \mathrm{P}_*]]\,\mathbf{p}\,[\mathrm{P}_* |\, [n, r]]\,\mathrm{d}P_*\,. \tag{5}$$

Table 1: Example calculations using Eq. (8), with $a = 1$ and $b = 1000$ in the prior $\mathbf{p}[\mathrm{P}_*]$.

n	r	P_f	n	r	P_f	n	r	P_f	n	r	P_f
10	1	0.02	100	1	0.16	200	1	0.27	1000	1	0.56
10	2	0.03	100	2	0.23	200	2	0.37	1000	2	0.70
10	3	0.04	100	3	0.30	200	3	0.46	1000	3	0.80
10	4	0.05	100	4	0.35	200	4	0.54	1000	4	0.87
10	5	0.06	100	5	0.41	200	5	0.60	1000	5	0.91
10	10	0.10	100	10	0.62	200	10	0.82	1000	10	0.99

Now, the issue of being able to count the number of 'successes' r_* from n experiments is addressed. The objective value of a particular optima may be evaluated, $\mathrm{F}_j\left[\mathbf{x}_*^j\right]$, and values amongst different optima may be compared: *e.g.* is $\mathrm{F}_c\left[\mathbf{x}_*^c\right]$ a better optima than $\mathrm{F}_d\left[\mathbf{x}_*^d\right]$? Hence, from n experiments, we can determine what was relatively the best optimum solution found $\boldsymbol{\chi}_*$, and we can count how many solutions ρ_* ('relative successes') correspond to this $\boldsymbol{\chi}_*$. If we assume that the probability of finding the global optimum from a random starting point P_*, is *at least* the probability of finding any other P_*^j

$$\mathrm{P}_* \geq \mathrm{P}_*^j \,\forall\, j \,, \tag{6}$$

then it follows that

$$\mathrm{P}\left[r_* \geq 1 | \left[n, r\right]\right] \geq \mathrm{P}\left[\rho_* \geq 1 | \left[n, r\right]\right] \,. \tag{7}$$

That is, the probability of having had at least one absolute 'success' r_*, after n experiments, is greater than the probability of having at least one 'relative success' ρ_*. This is property that may have to be justified in the optimization algorithm; *e.g.*, move-limits should be sufficiently large?

Substitution of (4) and (2) in (5), and, in turn, substitution of the binomial (3) and beta (1) distributions, permits straight-forward cancellations and simplification [1] to

$$\begin{aligned} \mathrm{P}_f = \mathrm{P}\left[r_* \geq 1 | \left[n, r\right]\right] \geq \mathrm{P}\left[\rho_* \geq 1 | \left[n, r\right]\right] = \\ 1 - \beta[r + a, 2n - r + b]/\beta[r + a, n - r + b] = \frac{\Gamma[n + a + b]\Gamma[2n - r + b]}{\Gamma[2n + a + b]\Gamma[n - r + b]} \,. \end{aligned} \tag{8}$$

Further cancellations are made in the factorials, and, in numerical terms, the calculation is dealt with as per the Python function given verbatim below.

In Table 1 some example calculations are given with a severely pessimistic estimate of $\mathbf{p}[\mathrm{P}_*]$, with $a = 1$ and $b = 1000$ (see Fig. 1).

We can see how the confidence in having found the global optimum P_f changes for different values of n and r. For example, with 10 experiments (n), even if we have a success rate of 1 in 2 ($r = 5$), our confidence barely rounds-up to 1 %. However, if we

do 100 experiments, and we have only between 5 and 10 successes, then the confidence of having found the global optimum increases to about 50%. This is not bad, given a prior expectation of $\frac{1}{1+1000} \approx \frac{1}{1000}$ (0.1 %). Given 200 experiments, and 10 successes, our confidence increases to a fairly comfortable 80%, *etc.*.

```
#
def P_f(ni,ri,a,b):
#
#   see Bolton (2004)
#
    tmp=1e0; abar=a+b-1; bbar=b-ri-1
    for i in range(1,ni+1):
        tmp=tmp*(ni+i+bbar)/(ni+i+abar)
#
    return 1e0-tmp
#
```

2 Applied to simple topology optimization

Fittingly avoiding the details of the optimization problem, the procedure is tested on the well-known Python implementation of the classical topology optimization problem [3]. The implementation is taken as is, the starting point is set to an array of values, with each individual element sampled from a uniform distribution between the lower and upper bounds of the problem, and the problem is executed 100 times (embarrassingly parallel). Throughout solutions will be characterized/compared based on objective function values and a measure of the number-of-holes' the topology has (a discrete version of a 'perimeter' quantity).

Based on a 1% threshold in the relative objective function values, the best solution—see Figure 2—is found $r = 33$ times. Based on a 5% difference in the perimeter quantity, the best solution was found $r = 37$ times. (We have been able to visually inspect whether these counts are representative.) Assuming a (horrible) Beta distribution prior with $a = 1$ and $b = 1000$, there is a 90% probability that we have found the global optimum[2].

For sake of interest; repeating the 100 experiments with a move-limit of 1.0 in the OC subsolver, results in $r = 60$, almost double the number of 'successes'; which brings the probability of having found the global optimum to 99%. That is, this manner of experimentation may also be a rigorous way of comparing the success of different optimization algorithms and parameter settings—while continuously building up confidence in the best known solution.

[2]Note that the particular filtering applied to the standard problem serves to explicitly restrict the number of solutions; we would have more candidate optima and fewer 'successes' for a smaller value of the filter[4].

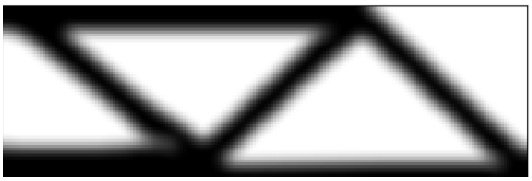

(a) Best solution found

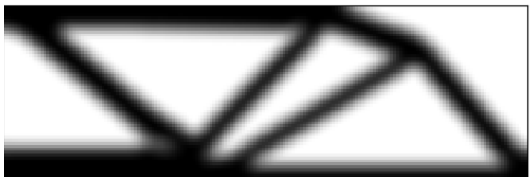

(b) Example of a local optima

Figure 2: Standard/default test problem; $n = 100$ experiments.

References

[1] J.A. Snyman and L.P. Fatti. A multi-start global minimization algorithm with dynamic search trajectories. *Journal of optimization theory and applications*, 54(1):121–141, 1987.

[2] F.M. Dekking, C. Kraaikamp, H. P. Lopuhaä, and L. E. Meester. *A modern introduction to probability and statistics: understanding why and how.* Springer Science & Business Media, 2005.

[3] Ole Sigmund. A 99 line topology optimization code written in matlab. *Structural and multidisciplinary optimization*, 21(2):120–127, 2001.

[4] Dirk Munro and Albert A Groenwold. On design-set restriction in sand topology optimization. *Structural and Multidisciplinary Optimization*, 57(4):1579–1592, 2018.